

PRD 17: Isomer Enumeration via Molecular Graph Theory

Pure Thought AI Challenge 17

Pure Thought AI Challenges Project

January 18, 2026

Abstract

This document presents a comprehensive Product Requirement Document (PRD) for implementing a pure-thought computational challenge. The problem can be tackled using only symbolic mathematics, exact arithmetic, and fresh code—no experimental data or materials databases required until final verification. All results must be accompanied by machine-checkable certificates.

Contents

Domain: Chemistry Combinatorics

Timeline: 4-6 months

Difficulty: Medium-High

Prerequisites: Graph theory, group theory (Pólya enumeration), combinatorial optimization, SAT solving

0.1 1. Problem Statement

0.1.1 Scientific Context

Isomers are molecules with the same chemical formula but different atomic arrangements. Enumerating all isomers for a given formula is a fundamental problem in chemistry:

Types of Isomers:

- **Structural isomers:** Different connectivity graphs (e.g., butane vs isobutane: both $\text{CH}_3\text{CH}_2\text{CH}_3$)
- **Stereoisomers:** Same connectivity, different 3D spatial arrangement
- **Conformers:** Same structure, different rotations around single bonds

Challenges:

- Number of isomers grows exponentially with molecular size
- Chemical valence rules constrain graphs (C forms 4 bonds, O forms 2, H forms 1)
- Symmetry: many graphs are equivalent up to atom relabeling (automorphisms)

Current Methods:

- **Brute force:** Generate all graphs, filter by valence—combinatorial explosion
- **Chemical databases:** Enumerate known structures—incomplete for large molecules
- **SMILES enumeration:** String-based, but misses many structures

Pure Thought Approach:

- Use **P 19th century mathematician György Pólya's enumeration theorem** to count distinct graphs
- Generate isomers systematically using **canonical labeling** (avoids duplicates)
- Apply chemical constraints as **SAT/SMT problems**
- Certify completeness: prove all isomers found

0.1.2 Core Question

Can we enumerate ALL isomers for a molecular formula $C_x H_y O_z \dots$ using ONLY graph theory and combinatorics?

Specifically:

- Given formula (e.g., CHO), generate all structurally distinct molecular graphs
- Apply valence constraints (C: 4, O: 2, N: 3, H: 1)
- Remove duplicate graphs via canonical labeling (nauty algorithm)
- Verify completeness: prove no isomers missed
- Extend to stereoisomers (chirality, E/Z isomerism)
- Export as SMILES strings + 3D geometries

0.1.3 Why This Matters

Theoretical Impact:

- Connects pure combinatorics to molecular chemistry
- Provides exact enumeration (no sampling or approximation)
- Algorithmic chemistry: automated structure generation

Practical Benefits:

- Drug discovery: enumerate all possible drug candidates with formula
- Materials design: explore chemical space systematically
- Retrosynthesis: identify alternative synthetic routes

Pure Thought Advantages:

- Valence rules are purely graph-theoretic
 - Pólya theory provides exact counts
 - No experimental data needed
 - Certificates of completeness via SAT solvers
-

0.2 2. Mathematical Formulation

0.2.1 Problem Definition

A molecular graph is $G = (V, E)$ where:

- V = vertices (atoms)
- E = edges (bonds), with multiplicities (single, double, triple)

Valence constraint: Each atom v has degree $\deg(v) = \text{valence}(\text{element}(v))$

- C: deg = 4
- O: deg = 2
- N: deg = 3
- H: deg = 1

Counting bond multiplicity:

```
1 deg(v) = len(neighbors(v)) bond_order(v, u)
```

Isomer Enumeration Problem:

```
1 Input: Molecular formula {n_C carbons, n_H hydrogens, n_O oxygens, ...}
2 Output: Set S of non-isomorphic connected molecular graphs satisfying
    valence constraints
```

Isomorphism: Two graphs G, G' are isomorphic if there exists bijection $: V \rightarrow V'$ preserving adjacency and atom types.

Certificate: For each isomer $G \in S$:

- Valence check: $v, \deg(v) = \text{valence}(v)$
- Canonicity: G is in canonical form (no other isomorphic graph generated)
- Completeness proof: All graphs in S are non-isomorphic + no missing isomers

0.2.2 Pólya Enumeration Theorem

For counting up to symmetry:

```
1 N = (1/|G|) * sum(g in G) cycle_index(g)
```

where G is the symmetry group, and cycle_i index counts fixed points under each symmetry.

0.2.3 Input/Output Specification

Input:

```
1 from typing import Dict
2 import networkx as nx
3
4 class MolecularFormula:
5     elements: Dict[str, int] # {'C': 4, 'H': 10} for C, H
6
7     # Optional constraints
8     allow_double_bonds: bool = True
9     allow_triple_bonds: bool = False
10    allow_rings: bool = True
11    max_ring_size: int = 8
```

Output:

```

1  class IsomerCertificate:
2      formula: MolecularFormula
3
4      # Enumerated isomers
5      isomers: List[nx.Graph]  # List of molecular graphs
6      num_isomers: int
7
8      # Canonical representations
9      canonical_smiles: List[str]  # SMILES strings
10     adjacency_matrices: List[np.ndarray]
11
12     # Verification
13     p_alya_count: int  # Theoretical count from Palya theorem
14     completeness_proof: str  # SAT/SMT certificate
15
16     # Statistics
17     num_with_rings: int
18     num_with_double_bonds: int
19     degree_distribution: Dict[int, int]
20
21     # Export
22     mol_files: List[Path]  # 3D structures (.mol, .xyz)
23     smiles_file: Path

```

0.3 3. Implementation Approach

0.3.1 Phase 1: Simple Enumeration for Small Molecules (Month 1)

Implement brute-force for validation:

```

1  import networkx as nx
2  from itertools import combinations
3  from typing import List
4
5  VALENCES = {'C': 4, 'H': 1, 'O': 2, 'N': 3, 'S': 2, 'P': 3, 'F': 1,
6      'Cl': 1}
7
8  def generate_all_graphs_brute_force(formula: MolecularFormula) ->
9      List[nx.Graph]:
10     """
11         Brute force: try all possible connectivity patterns.
12         Only works for very small molecules (    10 atoms).
13     """
14     # Create vertex list
15     atoms = []
16     for elem, count in formula.elements.items():
17         atoms.extend([elem] * count)
18
19     n = len(atoms)
20
21     # All possible edge sets (choose subset of n(n-1)/2 possible edges)
22     max_edges = n * (n-1) // 2

```

```

22     valid_graphs = []
23
24     # Iterate over all possible edge sets
25     all_possible_edges = list(combinations(range(n), 2))
26
27     for num_edges in range(n-1, max_edges+1):  # At least n-1 for
28         connectivity
29         for edge_set in combinations(all_possible_edges, num_edges):
30             G = nx.Graph()
31             G.add_nodes_from(range(n))
32
33             # Assign atom types
34             for i, atom in enumerate(atoms):
35                 G.nodes[i]['element'] = atom
36
37             # Add edges (all single bonds for now)
38             for (u, v) in edge_set:
39                 G.add_edge(u, v, bond_order=1)
40
41             # Check valence
42             if satisfies_valence(G) and nx.is_connected(G):
43                 valid_graphs.append(G)
44
45     # Remove isomorphic duplicates
46     unique_graphs = remove_isomorphic_duplicates(valid_graphs)
47
48     return unique_graphs
49
50 def satisfies_valence(G: nx.Graph) -> bool:
51     """Check if all atoms satisfy valence constraints."""
52     for node in G.nodes:
53         elem = G.nodes[node]['element']
54         required_valence = VALENCES[elem]
55
56         # Degree = sum of bond orders
57         degree = sum(G[node][nbr].get('bond_order', 1) for nbr in
58                     G.neighbors(node))
59
60         if degree != required_valence:
61             return False
62
63     return True
64
65 def remove_isomorphic_duplicates(graphs: List[nx.Graph]) ->
66     List[nx.Graph]:
67     """
68     Remove isomorphic graphs using nauty canonical labeling.
69     """
70     from pynauty import Graph as PynautyGraph, certificate
71
72     unique = []
73     seen_certificates = set()
74
75     for G in graphs:
76         # Convert to pynauty format
77         cert = compute_canonical_certificate(G)

```

```

75
76     if cert not in seen_certificates:
77         seen_certificates.add(cert)
78         unique.append(G)
79
80     return unique

```

Validation: Enumerate CH—should find 2 isomers (butane, isobutane).

0.3.2 Phase 2: Canonical Labeling with nauty (Months 1-2)

Use nauty algorithm for efficient isomorphism checking:

```

1 from pynauty import Graph as PynautyGraph, autgrp, certificate
2
3 def canonical_label_molecular_graph(G: nx.Graph) -> str:
4     """
5     Compute canonical labeling using nauty.
6
7     Returns string certificate uniquely identifying isomorphism class.
8     """
9     n = len(G.nodes)
10
11    # Partition vertices by atom type (nauty requires integer colors)
12    elem_to_color = {elem: i for i, elem in
13                     enumerate(set(VALENCES.keys()))}
14
15    coloring = [elem_to_color[G.nodes[v]['element']] for v in range(n)]
16
17    # Convert to pynauty format
18    adjacency = {v: list(G.neighbors(v)) for v in G.nodes}
19
20    pynauty_graph = PynautyGraph(
21        number_of_vertices=n,
22        directed=False,
23        adjacency_dict=adjacency,
24        vertex_coloring=[coloring]
25    )
26
27    # Compute canonical certificate
28    cert = certificate(pynauty_graph)
29
30    return str(cert)
31
32 def is_canonical(G: nx.Graph) -> bool:
33     """
34     Check if graph is in canonical form.
35
36     Canonical form: relabeling that is lexicographically smallest.
37     """
38     cert_original = canonical_label_molecular_graph(G)
39
40     # Try all permutations (expensive just for validation)
41     import itertools
42
43     n = len(G.nodes)

```

```
43     for perm in itertools.permutations(range(n)):
44         G_perm = relabel_graph(G, perm)
45         cert_perm = canonical_label_molecular_graph(G_perm)
46
47         if cert_perm < cert_original:
48             return False # Found smaller labeling
49
50     return True
```

0.3.3 Phase 3: Systematic Graph Generation (Months 2-4)

Use orderly generation (McKay's algorithm):

```

1 def orderly_generation(formula: MolecularFormula) -> List[nx.Graph]:
2     """
3         Generate graphs in canonical (orderly) manner.
4
5         Avoids generating isomorphic duplicates.
6
7         Based on McKay's orderly generation algorithm.
8         """
9
10    atoms = expand_formula(formula) # ['C', 'C', 'C', 'C', 'H', 'H',
11        ...]
12    n = len(atoms)
13
14    isomers = []
15
16    # Start with empty graph
17    G_init = nx.Graph()
18    G_init.add_nodes_from(range(n))
19    for i, elem in enumerate(atoms):
20        G_init.nodes[i]['element'] = elem
21
22    # Recursively add edges in canonical order
23    def generate_recursive(G, edge_candidates):
24        # Check if valid molecular graph
25        if is_complete_and_valid(G):
26            # Check canonical
27            if isCanonicalUnderAutomorphism(G):
28                isomers.append(G.copy())
29
30        # Pruning: stop if overvalent
31        if has_overvalent_atom(G):
32            return
33
34        # Add next edge (in canonical order)
35        for (u, v) in edge_candidates:
36            if can_add_edge(G, u, v):
37                G_new = G.copy()
38                G_new.add_edge(u, v, bond_order=1)
39
40                # Recursively expand
41                remaining_candidates = [(i, j) for (i, j) in
42                    edge_candidates
43                    if (i, j) > (u, v)]

```

```

42         generate_recursive(G_new, remaining_candidates)
43
44     # All possible edges
45     edge_candidates = list(combinations(range(n), 2))
46     generate_recursive(G_init, edge_candidates)
47
48     return isomers
49
50 def is_complete_and_valid(G: nx.Graph) -> bool:
51     """
52     Check if graph is a complete valid molecule.
53
54     - All atoms satisfy valence
55     - Graph is connected
56     """
57     if not nx.is_connected(G):
58         return False
59
60     for node in G.nodes:
61         elem = G.nodes[node]['element']
62         degree = G.degree(node)
63
64         if degree != VALENCES[elem]:
65             return False
66
67     return True

```

0.3.4 Phase 4: Pólya Enumeration (Months 4-5)

Count isomers using Pólya's theorem:

```

1 from sympy import symbols, expand, Poly
2 from sympy.combinatorics import PermutationGroup, Permutation
3
4 def polya_count_isomers(formula: MolecularFormula) -> int:
5     """
6     Use Pólya enumeration theorem to count non-isomorphic graphs.
7
8     This gives theoretical count doesn't enumerate structures.
9     """
10    atoms = expand_formula(formula)
11    n = len(atoms)
12
13    # Symmetry group: permutations preserving atom types
14    # E.g., for C H : permutations of 4 C's permutations of
15    # 10 H's
16
17    C_indices = [i for i, a in enumerate(atoms) if a == 'C']
18    H_indices = [i for i, a in enumerate(atoms) if a == 'H']
19
20    # Generate symmetric group on each atom type
21    perms_C = PermutationGroup([Permutation(C_indices)])
22    perms_H = PermutationGroup([Permutation(H_indices)])
23
24    # Full group: product
25    G = combine_permutation_groups(perms_C, perms_H)

```

```

25
26     # Cycle index polynomial
27     x = symbols(f'x0:{n*(n-1)//2}') # Variables for each edge
28
29     cycle_poly = compute_cycle_index(G, x)
30
31     # Substitute x_i      1 + t (count graphs with/without each edge)
32     t = symbols('t')
33     cycle_poly_sub = cycle_poly.subs({xi: 1+t for xi in x})
34
35     # Extract coefficient of t^m where m = number of edges
36     # For tree: m = n-1
37     # For general graphs with cycles: various m
38
39     poly_expanded = expand(cycle_poly_sub)
40     coeffs = Poly(poly_expanded, t).all_coeffs()
41
42     # Number of isomers with m edges
43     isomer_counts = {m: coeffs[m] for m in range(len(coeffs))}

44
45     # Filter for chemically valid (satisfies valence)
46     # This is approximate exact filtering requires enumeration
47
48     total_isomers = sum(isomer_counts.values())
49
50     return total_isomers
51
52 def compute_cycle_index(G: PermutationGroup, variables: List) -> Poly:
53     """
54         Compute cycle index polynomial for permutation group G.
55
56         Z(G) = (1/|G|) * sum_{g in G} prod_i x_i^{c_i(g)}
57
58         where c_i(g) = number of i-cycles in permutation g.
59     """
60     from sympy import Rational
61
62     cycle_poly = 0
63
64     for g in G.generate():
65         # Cycle structure of permutation g
66         cycles = g.cyclic_form
67
68         # Product over cycle lengths
69         term = 1
70         for cycle in cycles:
71             cycle_len = len(cycle)
72             term *= variables[cycle_len - 1]
73
74         cycle_poly += term
75
76     cycle_poly /= len(list(G.generate()))
77
78     return cycle_poly

```

0.3.5 Phase 5: Stereoisomer Enumeration (Months 5-6)

Extend to 3D stereochemistry:

```

1 def enumerate_stereoisomers(molecular_graph: nx.Graph) -> List:
2     """
3         For each structural isomer, enumerate stereoisomers.
4
5         - Chiral centers: tetrahedral carbons with 4 different substituents
6         - E/Z isomers: double bonds with different substituents
7         - Conformers: rotations around single bonds (separate problem)
8     """
9     stereoisomers = []
10
11     # Find chiral centers
12     chiral_centers = find_chiral_carbons(molecular_graph)
13
14     # 2^n stereoisomers for n chiral centers (R/S configurations)
15     for config in itertools.product(['R', 'S'],
16                                     repeat=len(chiral_centers)):
17         G_stereo = molecular_graph.copy()
18
19         for center, chirality in zip(chiral_centers, config):
20             G_stereo.nodes[center]['chirality'] = chirality
21
22         stereoisomers.append(G_stereo)
23
24     # E/Z isomers (double bonds)
25     double_bonds = [(u, v) for u, v in molecular_graph.edges
26                      if molecular_graph[u][v]['bond_order'] == 2]
27
28     for bond in double_bonds:
29         # Check if E/Z isomerism possible
30         if has_EZ_isomerism(molecular_graph, bond):
31             # Generate both E and Z forms
32             # ... (geometric isomer generation)
33             pass
34
35     return stereoisomers
36
37 def find_chiral_carbons(G: nx.Graph) -> List[int]:
38     """
39         Identify chiral centers (sp   carbons with 4 different groups).
40     """
41     chiral = []
42
43     for node in G.nodes:
44         if G.nodes[node]['element'] != 'C':
45             continue
46
47         if G.degree(node) != 4:
48             continue # Must be tetrahedral
49
50         # Check if 4 neighbors are distinct
51         neighbors = list(G.neighbors(node))
52         if all_distinct_substituents(G, neighbors):
53             chiral.append(node)

```

```

53
54     return chiral

```

0.3.6 Phase 6: Export and Validation (Month 6)

Generate output formats (SMILES, 3D structures):

```

1  from rdkit import Chem
2  from rdkit.Chem import AllChem
3
4  def export_isomers(isomers: List[nx.Graph], output_dir: Path):
5      """
6          Export isomers as SMILES and 3D structures.
7      """
8      smiles_list = []
9
10     for i, G in enumerate(isomers):
11         # Convert to RDKit molecule
12         mol = nx_graph_to_rdkit(G)
13
14         # Generate SMILES
15         smiles = Chem.MolToSmiles(mol)
16         smiles_list.append(smiles)
17
18         # Generate 3D coordinates (force field optimization)
19         mol_3d = Chem.AddHs(mol)
20         AllChem.EmbedMolecule(mol_3d)
21         AllChem.UFFOptimizeMolecule(mol_3d)
22
23         # Save as MOL file
24         Chem.MolToMolFile(mol_3d, str(output_dir /
25             f'isomer_{i:03d}.mol'))
26
27     # Save all SMILES
28     with open(output_dir / 'isomers.smi', 'w') as f:
29         for smi in smiles_list:
30             f.write(smi + '\n')
31
32     def nx_graph_to_rdkit(G: nx.Graph) -> Chem.Mol:
33         """
34             Convert NetworkX molecular graph to RDKit Mol object.
35         """
36         mol = Chem.RWMol()
37
38         # Add atoms
39         node_to_idx = {}
40         for node in G.nodes:
41             elem = G.nodes[node]['element']
42             atom = Chem.Atom(elem)
43             idx = mol.AddAtom(atom)
44             node_to_idx[node] = idx
45
46         # Add bonds
47         for u, v in G.edges:
48             bond_order = G[u][v].get('bond_order', 1)

```

```

49     if bond_order == 1:
50         bond_type = Chem.BondType.SINGLE
51     elif bond_order == 2:
52         bond_type = Chem.BondType.DOUBLE
53     elif bond_order == 3:
54         bond_type = Chem.BondType.TRIPLE
55
56     mol.AddBond(node_to_idx[u], node_to_idx[v], bond_type)
57
58 return mol.GetMol()

```

0.4 4. Example Starting Prompt

```

1 You are a computational chemist implementing isomer enumeration via
2 graph theory. Generate ALL
3 structural isomers for molecular formulas using ONLY combinatorial
4 algorithms no databases.
5
6 OBJECTIVE: Enumerate all C HO isomers, verify completeness,
7 export as SMILES.
8
9 PHASE 1 (Month 1): Brute force baseline
10 - Implement basic graph generation for C H
11 - Apply valence constraints (C:4, H:1, O:2)
12 - Remove duplicates using nauty canonical labeling
13 - Verify: find exactly 2 isomers (butane, isobutane)
14
15 PHASE 2 (Months 1-2): Canonical labeling
16 - Implement nauty algorithm for molecular graphs
17 - Color vertices by element type
18 - Test isomorphism detection on 100 random graph pairs
19
20 PHASE 3 (Months 2-4): Orderly generation
21 - Implement McKay's orderly algorithm
22 - Generate graphs in canonical order (avoids duplicates)
23 - Test on C H : should find 3 isomers
24
25 PHASE 4 (Months 4-5): P lya enumeration
26 - Compute symmetry group for C HO
27 - Calculate cycle index polynomial
28 - Compare P lya count to generated count (must match!)
29
30 PHASE 5 (Months 5-6): Stereoisomers
31 - Identify chiral centers in each structural isomer
32 - Enumerate R/S configurations
33 - Handle E/Z isomerism for double bonds
34
35 PHASE 6 (Month 6): Export and validation
36 - Convert all isomers to SMILES strings
37 - Generate 3D geometries using RDKit
38 - Cross-check against PubChem database (for validation only)
39
40 SUCCESS CRITERIA:

```

```
38 - MVR: C H and C H correctly enumerated
39 - Strong: C HO complete enumeration, all unique structures
40 - Publication: Systematic study up to C , comparison to P ly a counts
41
42 VERIFICATION:
43 - Generated count matches P ly a theoretical count
44 - All SMILES strings valid (parseable by RDKit)
45 - No duplicate structures (canonical checking)
46 - Cross-reference with PubChem (should find all known isomers)
47
48 Pure graph theory + combinatorics. No chemical databases until final
   validation.
49 All results certificate-based with completeness proofs.
```

0.5 5. Success Criteria

MVR (2 months): CH, CH correct, nauty working

Strong (4-5 months): CHO complete, Pólya counts verified

Publication (6 months): Systematic enumeration up to C, stereoisomers included

0.6 6. Verification Protocol

- Compare generated counts to Pólya theoretical values
 - Cross-check SMILES against PubChem
 - Validate 3D geometries with quantum chemistry (DFT single-points)
 - Verify canonical labeling (no duplicates)
-

0.7 7. Resources Milestones

References:

- McKay (1998): "Isomorph-Free Exhaustive Generation"
- Pólya (1937): "Kombinatorische Anzahlbestimmungen"
- Read Corneil (1977): "Graph Isomorphism Algorithms"

Milestones:

- Month 2: nauty integration complete
 - Month 4: Orderly generation working
 - Month 6: Full CHO enumeration + stereoisomers
-

0.8 8. Extensions

- **Reactivity Prediction:** Which isomers are most stable/reactive?
 - **Retrosynthesis:** Enumerate synthetic routes
 - **Protein Folding:** Graph enumeration for polymer conformations
-

End of PRD 17